

Machine Learning - Exercise 2: Logistic Regression

Instructions

This file contains code that helps you get started on the logistic regression exercise. You will need to complete the following functions in this exercise:

- % **plotData.m**
- % **sigmoid.m**
- % **costFunction.m**
- % **predict.m**

```
%% Initialization
clear ; close all; clc

%% Load Data
% The first two columns contains the exam scores and the third column
% contains the label.

data = load('ex2data1.txt');
X = data(:, [1, 2]); y = data(:, 3);

%% ===== Part 1: Plotting =====
% We start the exercise by first plotting the data to understand the
% the problem we are working with.

fprintf(['Plotting data with + indicating (y = 1) examples and o ' ...
        'indicating (y = 0) examples.\n']);
```


Plotting data with + indicating (y = 1) examples and o indicating (y = 0) examples.

```
%%=====Finish plotData Below=====
plotData(X, y);

% Put some labels
hold on;
% Labels and Legend
xlabel('Exam 1 score')
ylabel('Exam 2 score')

% Specified in plot order
legend('Admitted', 'Not admitted')
```

```
hold off;
```



```
fprintf('\nProgram paused. Press enter to continue.\n');
```

Program paused. Press enter to continue.

```
pause;
```

```
%% ===== Part 2: Compute Cost and Gradient =====
```

```
% In this part of the exercise, you will implement the cost and gradient  
% for logistic regression. You need to complete the code in  
% costFunction.m
```

```
% Setup the data matrix appropriately, and add ones for the intercept term  
[m, n] = size(X);
```

```
% Add intercept term to x and X_test  
X = [ones(m, 1) X];
```

```
% Initialize fitting parameters  
initial_theta = zeros(n + 1, 1);
```

```
% =====Finish costFunction to Compute and display initial cost and  
% gradient=====
```

```
[cost, grad] = costFunction(initial_theta, X, y);
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1
    -0.1000
         0
         0
grad = 3x1
    -0.1000
   -12.0092
         0
grad = 3x1
    -0.1000
   -12.0092
   -11.2628
```

```
fprintf('Cost at initial theta (zeros): %f\n', cost);
```

Cost at initial theta (zeros): 0.693147

```
fprintf('Expected cost (approx): 0.693\n');
```

Expected cost (approx): 0.693

```
fprintf('Gradient at initial theta (zeros): \n');
```

Gradient at initial theta (zeros):

```
fprintf(' %f \n', grad);
```

```
-0.100000
-12.009217
-11.262842
```

```
fprintf('Expected gradients (approx):\n -0.1000\n -12.0092\n -11.2628\n');
```

Expected gradients (approx):

```
-0.1000
-12.0092
-11.2628
```

% Compute and display cost and gradient with non-zero theta

```
test_theta = [-24; 0.2; 0.2];
[cost, grad] = costFunction(test_theta, X, y);
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1
    0.0429
         0
         0
grad = 3x1
    0.0429
    2.5662
         0
grad = 3x1
    0.0429
    2.5662
    2.6468
```

```
fprintf('\nCost at test theta: %f\n', cost);
```

Cost at test theta: 0.218330

```
fprintf('Expected cost (approx): 0.218\n');
```

Expected cost (approx): 0.218

```
fprintf('Gradient at test theta: \n');
```

Gradient at test theta:

```
fprintf(' %f \n', grad);
```

0.042903
2.566234
2.646797

```
fprintf('Expected gradients (approx):\n 0.043\n 2.566\n 2.647\n');
```

Expected gradients (approx):

0.043
2.566
2.647

```
fprintf('\nProgram paused. Press enter to continue.\n');
```

Program paused. Press enter to continue.

```
pause;
```

```
%% ===== Part 3: Optimizing using fminunc =====  
% In this exercise, you will use a built-in function (fminunc) to find the  
% optimal parameters theta.  
  
% Set options for fminunc  
options = optimset('GradObj', 'on', 'MaxIter', 400);  
  
% Run fminunc to obtain the optimal theta  
% This function will return theta and the cost  
[theta, cost] = ...  
    fminunc(@(t)(costFunction(t, X, y)), initial_theta, options);
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1  
    -0.1000  
         0  
         0  
  
grad = 3x1  
    -0.1000  
   -12.0092  
         0  
  
grad = 3x1  
    -0.1000  
   -12.0092  
   -11.2628
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1
  0.4000
   0
   0
```

```
grad = 3x1
  0.4000
 20.8129
   0
```

```
grad = 3x1
  0.4000
 20.8129
 21.8482
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1
  0.4000
   0
   0
```

```
grad = 3x1
  0.4000
 20.8129
   0
```

```
grad = 3x1
  0.4000
 20.8129
 21.8482
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1
  0.3757
   0
   0
```

```
grad = 3x1
  0.3757
 19.4850
   0
```

```
grad = 3x1
  0.3757
 19.4850
 20.4520
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1
  0.0799
   0
   0
```

```
grad = 3x1
  0.0799
  0.2760
   0
```

```
grad = 3x1
  0.0799
  0.2760
  1.0541
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3x1
  0.0712
   0
   0
```

```
grad = 3x1
  0.0712
 -0.2986
   0
```

```
grad = 3x1
  0.0712
```

```

-0.2986
0.4482
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3x1
0.0671
0
0
grad = 3x1
0.0671
-0.5505
0
grad = 3x1
0.0671
-0.5505
0.1437
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3x1
0.0613
0
0
grad = 3x1
0.0613
-0.8535
0
grad = 3x1
0.0613
-0.8535
-0.3370
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3x1
0.0594
0
0
grad = 3x1
0.0594
-0.8391
0
grad = 3x1
0.0594
-0.8391
-0.5966
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3x1
0.0636
0
0
grad = 3x1
0.0636
-0.4306
0
grad = 3x1
0.0636
-0.4306
-0.4168
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3x1
0.0684
0
0
grad = 3x1
0.0684
-0.0797
0

```

```

grad = 3×1
  0.0684
 -0.0797
 -0.1143
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
  0.0698
   0
   0
grad = 3×1
  0.0698
  0.0119
   0
grad = 3×1
  0.0698
  0.0119
 -0.0090
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
  0.0703
   0
   0
grad = 3×1
  0.0703
  0.0352
   0
grad = 3×1
  0.0703
  0.0352
  0.0231
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
  0.0712
   0
   0
grad = 3×1
  0.0712
  0.0873
   0
grad = 3×1
  0.0712
  0.0873
  0.0975
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
  0.0725
   0
   0
grad = 3×1
  0.0725
  0.1583
   0
grad = 3×1
  0.0725
  0.1583
  0.1997
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
  0.0747
   0
   0
grad = 3×1
  0.0747
  0.2798

```

```

0
grad = 3×1
0.0747
0.2798
0.3749
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0781
0
0
grad = 3×1
0.0781
0.4708
0
grad = 3×1
0.0781
0.4708
0.6501
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0836
0
0
grad = 3×1
0.0836
0.7790
0
grad = 3×1
0.0836
0.7790
1.0945
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0921
0
0
grad = 3×1
0.0921
1.2650
0
grad = 3×1
0.0921
1.2650
1.7954
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.1049
0
0
grad = 3×1
0.1049
2.0159
0
grad = 3×1
0.1049
2.0159
2.8796
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.1229
0
0
grad = 3×1
0.1229

```



```

3.1076
0
grad = 3×1
0.1229
3.1076
4.4621
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.1438
0
0
grad = 3×1
0.1438
4.4753
0
grad = 3×1
0.1438
4.4753
6.4797
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.1565
0
0
grad = 3×1
0.1565
5.5891
0
grad = 3×1
0.1565
5.5891
8.2727
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.1433
0
0
grad = 3×1
0.1433
5.5020
0
grad = 3×1
0.1433
5.5020
8.5397
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.1110
0
0
grad = 3×1
0.1110
4.5069
0
grad = 3×1
0.1110
4.5069
7.1859
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0753
0
0
grad = 3×1

```

```

0.0753
3.3358
0
grad = 3×1
0.0753
3.3358
5.0410
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0426
0
0
grad = 3×1
0.0426
2.1096
0
grad = 3×1
0.0426
2.1096
2.8361
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0223
0
0
grad = 3×1
0.0223
1.2304
0
grad = 3×1
0.0223
1.2304
1.4509
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0106
0
0
grad = 3×1
0.0106
0.6568
0
grad = 3×1
0.0106
0.6568
0.6686
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0051
0
0
grad = 3×1
0.0051
0.3497
0
grad = 3×1
0.0051
0.3497
0.3190
Warning: Colon operands must be real scalars. This warning will become an error in a future release.
grad = 3×1
0.0025
0
0

```

```
grad = 3×1
  0.0025
  0.1837
  0
```

```
grad = 3×1
  0.0025
  0.1837
  0.1676
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1
1.0e-03 *
  0.4029
  0
  0
```

```
grad = 3×1
  0.0004
  0.0315
  0
```

```
grad = 3×1
  0.0004
  0.0315
  0.0523
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1
1.0e-03 *
 -0.4449
  0
  0
```

```
grad = 3×1
 -0.0004
 -0.0418
  0
```

```
grad = 3×1
 -0.0004
 -0.0418
  0.0076
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1
1.0e-03 *
 -0.3904
  0
  0
```

```
grad = 3×1
 -0.0004
 -0.0464
  0
```

```
grad = 3×1
 -0.0004
 -0.0464
  0.0056
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1
1.0e-04 *
 -0.4409
  0
  0
```

```
grad = 3×1
 -0.0000
 -0.0179
  0
```

```
grad = 3×1
 -0.0000
 -0.0179
  0.0123
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1  
1.0e-04 *  
    0.8716  
    0  
    0
```

```
grad = 3×1  
1.0e-03 *  
    0.0872  
    0.9216  
    0
```

```
grad = 3×1  
    0.0001  
    0.0009  
    0.0086
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1  
1.0e-04 *  
    0.3714  
    0  
    0
```

```
grad = 3×1  
    0.0000  
    0.0019  
    0
```

```
grad = 3×1  
    0.0000  
    0.0019  
    0.0024
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1  
1.0e-05 *  
    0.5702  
    0  
    0
```

```
grad = 3×1  
1.0e-03 *  
    0.0057  
    0.3852  
    0
```

```
grad = 3×1  
1.0e-03 *  
    0.0057  
    0.3852  
    0.3048
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1  
1.0e-06 *  
    0.4374  
    0  
    0
```

```
grad = 3×1  
1.0e-04 *  
    0.0044  
    0.3340  
    0
```

```
grad = 3×1  
1.0e-04 *  
    0.0044  
    0.3340  
    0.2299
```

Warning: Colon operands must be real scalars. This warning will become an error in a future release.

```
grad = 3×1  
1.0e-07 *
```

```

    0.1535
    0
    0
grad = 3×1
1.0e-05 *
    0.0015
    0.1203
    0
grad = 3×1
1.0e-05 *
    0.0015
    0.1203
    0.1050
Local minimum found.

```

Optimization completed because the size of the gradient is less than the value of the optimality tolerance.

<stopping criteria details>

```

% fminunc finds costfunction's minimizer

% BONUS – use gradient descent you wrote from last HW to verify the theta
% results

% Print theta to screen
fprintf('Cost at theta found by fminunc: %f\n', cost);

```

Cost at theta found by fminunc: 0.203498

```
fprintf('Expected cost (approx): 0.203\n');
```

Expected cost (approx): 0.203

```
fprintf('theta: \n');
```

theta:

```
fprintf(' %f \n', theta);
```

```

-25.161343
0.206232
0.201472

```

```
fprintf('Expected theta (approx):\n');
```

Expected theta (approx):

```
fprintf(' -25.161\n 0.206\n 0.201\n');
```

```

-25.161
0.206
0.201

```

```

% Plot Boundary
plotDecisionBoundary(theta, X, y);


```

```

% Put some labels
hold on;
% Labels and Legend
xlabel('Exam 1 score')
ylabel('Exam 2 score')

% Specified in plot order
legend('Admitted', 'Not admitted')
hold off;

```



```
fprintf('\nProgram paused. Press enter to continue.\n');
```

Program paused. Press enter to continue.

```
pause;
```

```

%% ===== Part 4: Predict and Accuracies =====
% After learning the parameters, you'll like to use it to predict the
% outcomes
% on unseen data. In this part, you will use the logistic regression model
% to predict the probability that a student with score 45 on exam 1 and
% score 85 on exam 2 will be admitted.
%
% Furthermore, you will compute the training and test set accuracies of
% our model.
%

```

```
% Your task is to complete the code in predict.m

% Predict probability for a student with score 45 on exam 1
% and score 85 on exam 2

prob = sigmoid([1 45 85] * theta);
fprintf(['For a student with scores 45 and 85, we predict an admission ' ...
        'probability of %f\n'], prob);
```

For a student with scores 45 and 85, we predict an admission probability of 0.776291

```
fprintf('Expected value: 0.775 +/- 0.002\n\n');
```

Expected value: 0.775 +/- 0.002

```
% =====Finish predict to Compute accuracy on our training set=====
p = predict(theta, X);

fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);
```

Train Accuracy: 89.000000

```
fprintf('Expected accuracy (approx): 89.0\n');
```

Expected accuracy (approx): 89.0

```
fprintf('\n');
```

```
function plotData(X, y)
%PLOTDATA Plots the data points X and y into a new figure
% PLOTDATA(x,y) plots the data points with + for the positive examples
% and o for the negative examples. X is assumed to be a Mx2 matrix.

% Create New Figure
figure; hold on;

% ===== YOUR CODE HERE =====
% Instructions: Plot the positive and negative examples on a
%               2D plot, using the option 'k+' for the positive
%               examples and 'ko' for the negative examples.
%
% Find Indices of Positive and Negative Examples
pos = find(y==1); %pos = find(y)
neg = find(y==0); %neg = find(~y)

% pos = [];
% neg = [];
%
```

```

% for i = 1:length(y)
%     if y(i)==1
%         pos = [pos i];
%     elseif y(i) == 0
%         neg = [neg i];
%     end
% end

% Plot Examples
plot(X(pos, 1), X(pos, 2), 'k+', 'LineWidth', 2, ...
'MarkerSize', 7);
plot(X(neg, 1), X(neg, 2), 'ko', 'MarkerFaceColor', 'y', ...
'MarkerSize', 7);

% =====

hold off;

end

```

sigmoid/logistic function $g(z) = \frac{1}{1 + e^{-z}}$.

```

function g = sigmoid(z)
%SIGMOID Compute sigmoid function
%   g = SIGMOID(z) computes the sigmoid of z.

% You need to return the following variables correctly
g = zeros(size(z));

% ===== YOUR CODE HERE =====
% Instructions: Compute the sigmoid of each value of z (z can be a matrix,
%               vector or scalar).

g = 1./(1+exp(-z)); %g = 1./(1+e.^(-z))

% =====

end

```


$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] ,$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$h_{\theta}(x) = g(X\theta), \quad g(z) = \frac{1}{1 + e^{-z}}.$$

```
function [J, grad] = costFunction(theta, X, y)
%COSTFUNCTION Compute cost and gradient for logistic regression
% J = COSTFUNCTION(theta, X, y) computes the cost of using theta as the
% parameter for logistic regression and the gradient of the cost
% w.r.t. to the parameters.

% Initialize some useful values
m = length(y); % number of training examples

% You need to return the following variables correctly
J = 0;
grad = zeros(size(theta));

% ===== YOUR CODE HERE =====
% Instructions: Compute the cost of a particular choice of theta.
%               You should set J to the cost.
%               Compute the partial derivatives and set grad to the partial
%               derivatives of the cost w.r.t. each parameter in theta
%
% Note: grad should have the same dimensions as theta
%

h = sigmoid(X*theta); %100 by 1 vector

J = (1/m)*(-y'*log(h)-(1-y)'*log(1-h));

for j = 1:size(theta)
    grad(j) = (1/m)*(h-y)'*X(:,j)
end

% bonus: can you compute grad in one line?

% =====

end
```

$h_{\theta}(x) \geq 0.5$, predict 1, otherwise, predict 0

```
function p = predict(theta, X)
%PREDICT Predict whether the label is 0 or 1 using learned logistic
%regression parameters theta

m = size(X, 1); % Number of training examples

% You need to return the following variables correctly
p = zeros(m, 1);

% ===== YOUR CODE HERE =====
% Instructions: Complete the following code to make predictions using
%               your learned logistic regression parameters.
%               You should set p to a vector of 0's and 1's
%

h = sigmoid(X*theta);

p(h>=0.5) = 1;

% p = round(h)

% if h_theta >= 0.5
%     p = 1;
% else
%     p = 0;

% =====

end
```

Appendix [DO NOT CHANGE THESE]

```
function plotDecisionBoundary(theta, X, y)
%PLOTDECISIONBOUNDARY Plots the data points X and y into a new figure with
%the decision boundary defined by theta
% PLOTDECISIONBOUNDARY(theta, X,y) plots the data points with + for the
% positive examples and o for the negative examples. X is assumed to be
% a either
% 1) Mx3 matrix, where the first column is an all-ones column for the
%    intercept.
% 2) MxN, N>3 matrix, where the first column is all-ones

% Plot Data
plotData(X(:,2:3), y);
hold on
```

```

if size(X, 2) <= 3
    % Only need 2 points to define a line, so choose two endpoints
    plot_x = [min(X(:,2))-2, max(X(:,2))+2];

    % Calculate the decision boundary line
    plot_y = (-1./theta(3)).*(theta(2).*plot_x + theta(1));

    % Plot, and adjust axes for better viewing
    plot(plot_x, plot_y)

    % Legend, specific for the exercise
    legend('Admitted', 'Not admitted', 'Decision Boundary')
    axis([30, 100, 30, 100])
else
    % Here is the grid range
    u = linspace(-1, 1.5, 50);
    v = linspace(-1, 1.5, 50);

    z = zeros(length(u), length(v));
    % Evaluate z = theta*x over the grid
    for i = 1:length(u)
        for j = 1:length(v)
            z(i,j) = mapFeature(u(i), v(j))*theta;
        end
    end
    z = z'; % important to transpose z before calling contour

    % Plot z = 0
    % Notice you need to specify the range [0, 0]
    contour(u, v, z, [0, 0], 'LineWidth', 2)
end
hold off

end

function out = mapFeature(X1, X2)
% MAPFEATURE Feature mapping function to polynomial features
%
% MAPFEATURE(X1, X2) maps the two input features
% to quadratic features used in the regularization exercise.
%
% Returns a new feature array with more features, comprising of
% X1, X2, X1.^2, X2.^2, X1*X2, X1*X2.^2, etc..
%
% Inputs X1, X2 must be the same size
%
degree = 6;

```

```
out = ones(size(X1(:,1)));  
for i = 1:degree  
    for j = 0:i  
        out(:, end+1) = (X1.^(i-j)).*(X2.^j);  
    end  
end  
end
```